

# Predictive Modelling Project

SUBMITTED BY: RISHAB SINGLA

|

# Table of Contents

List Of tables .....	2
Table of figures .....	2
Problem 1: Linear Regression .....	5
Question 1.1. ....	6
Answer 1.1 .....	6
Question 1.2 .....	15
Answer 1.2 .....	15
Question 1.3 .....	16
Answer 1.3 .....	17
Question 1.4 .....	26
Answer 1.4 .....	26
Problem 2: Logistic Regression and LDA.....	28
Question 2.1 .....	28
Answer 2.1 .....	28
Question 2.2 .....	34
Answer 2.2 .....	34
Question 2.3 .....	36
Answer 2.3 .....	36
Question 2.4 .....	46
Answer 2.4 .....	46

## List Of tables

Table 1:Model comparison .....	26
Table 2:Model comparison .....	40
Table 3:Model comparison .....	45

## Table of figures

Figure 1: Dataset Information .....	6
Figure 2: Dataset Description.....	6
Figure 3: Null Values .....	7
Figure 4: Records with Null Values .....	7
Figure 5: Zero Value Records .....	7
Figure 6: Duplicated Records .....	8
Figure 7: Variable 'depth' .....	9
Figure 8: Variable 'carat' .....	9
Figure 9: Variable 'table' .....	9
Figure 10: Variable 'x' .....	10
Figure 11: Variable 'y' .....	10
Figure 12: Variable 'Z' .....	10
Figure 13: Variable 'Price' .....	11
Figure 14: Boxplot .....	11
Figure 15: Boxplot .....	11
Figure 16: Boxplot .....	12
Figure 17: Heatmap.....	12
Figure 18: Pairplot.....	13
Figure 19: Columns after Label Encoding .....	14
Figure 20: Boxplot After Outlier Treatment.....	14
Figure 21: Null Values .....	15
Figure 22: Null Values after Data Modification .....	15
Figure 23: Null values After Imputation.....	16
Figure 24: One Hot Encoding .....	17
Figure 25: Model 1 Regression Analysis.....	18
Figure 26: Scatter Plot Between Actual and Predicted Price.....	19
Figure 27: Variance Inflation Factor.....	19
Figure 28: Model 2 Regression Analysis.....	20
Figure 29: Scatter Plot Between Actual and Predicted Price.....	20
Figure 30: Variance Inflation Factor.....	21
Figure 31: Model 3 Regression Analysis.....	22
Figure 32: Scatter Plot Between Actual and Predicted Price.....	22
Figure 33: Variance Inflation Factor.....	23
Figure 34: Model 4 Regression Analysis.....	24

Figure 35: Variance Inflation Factor.....	24
Figure 36: Modified Columns.....	25
Figure 37: Model 5 Regression Analysis.....	25
Figure 38: Linear Regression Equation.....	26
Figure 39: Dataset Information.....	28
Figure 40: Dataset Description.....	29
Figure 41: Null Values .....	29
Figure 42: Holliday_Package Value Count .....	29
Figure 43: Foreign Value Count.....	29
Figure 44: Variable 'Salary'.....	30
Figure 45: Variable 'Age' .....	30
Figure 46: Variable 'Educ'.....	30
Figure 47: Variable 'No_young_children' .....	31
Figure 48; Variable 'No_older_children' .....	31
Figure 49: Count plot of Educ .....	32
Figure 50: Count Plot of no_young_children.....	32
Figure 51: Count plot of no_older_children .....	32
Figure 52: Count Plot of foreign.....	33
Figure 53: Heatmap.....	33
Figure 54: Pairplot.....	34
Figure 55: Label Encoding .....	35
Figure 56:Grid Search CV .....	35
Figure 57: Best Parameters.....	35
Figure 58: Linear Discriminant Analysis .....	36
Figure 59: ROC Curve .....	36
Figure 60: Classification Report .....	37
Figure 61: Confusion Matrix.....	37
Figure 62: ROC Curve .....	37
Figure 63: Classification Report .....	38
Figure 64: Confusion Matrix.....	38
Figure 65: ROC Curve .....	38
Figure 66: Classification Report .....	39
Figure 67: Confusion Matrix.....	39
Figure 68: ROC Curve .....	39
Figure 69: Classification Report .....	39
Figure 70: Confusion Matrix.....	40
Figure 71: Boxplot.....	41
Figure 72: Grid Search CV .....	41
Figure 73: Best Parameters.....	41
Figure 74: ROC Curve .....	42
Figure 75: Classification Report .....	42
Figure 76: Confusion Matrix.....	42

Figure 77: ROC Curve .....	43
Figure 78: Classification Report .....	43
Figure 79: Confusion Matrix.....	43
Figure 80: ROC Curve .....	44
Figure 81: Classification Report .....	44
Figure 82: Confusion Matrix.....	44
Figure 83: ROC Curve .....	45
Figure 84: Classification Report .....	45
Figure 85: ConfusionMatrix .....	45
Figure 86: Coefficients .....	46

### Problem 1: Linear Regression

You are hired by a company Gem Stones co ltd, which is a cubic zirconia manufacturer. You are provided with the dataset containing the prices and other attributes of almost 27,000 cubic zirconia (which is an inexpensive diamond alternative with many of the same qualities as a diamond). The company is earning different profits on different prize slots. You have to help the company in predicting the price for the stone on the bases of the details given in the dataset so it can distinguish between higher profitable stones and lower profitable stones so as to have better profit share. Also, provide them with the best 5 attributes that are most important.

Variable Name	Description
Carat	Carat weight of the cubic zirconia.
Cut	Describe the cut quality of the cubic zirconia. Quality is increasing order Fair, Good, Very Good, Premium, Ideal.
Color	Color of the cubic zirconia. With D being the worst and J the best.
Clarity	Clarity refers to the absence of the Inclusions and Blemishes. (In order from Worst to Best in terms of avg price) IF, VVS1, VVS2, VS1, VS2, SI1, SI2, I1
Depth	The Height of cubic zirconia, measured from the Culet to the table, divided by its average Girdle Diameter.
Table	The Width of the cubic zirconia's Table expressed as a Percentage of its Average Diameter.
Price	the Price of the cubic zirconia.
X	Length of the cubic zirconia in mm.
Y	Width of the cubic zirconia in mm.
Z	Height of the cubic zirconia in mm.

### Question 1.1.

Read the data and do exploratory data analysis. Describe the data briefly. (Check the null values, Data types, shape, EDA, duplicate values). Perform Univariate and Bivariate Analysis.

### Answer 1.1

The dataset consists of 26967 rows and 11 columns. 3 columns are of datatype object namely cut, clarity and color. All other columns are of continuous nature and datatype as int or float.

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 26967 entries, 0 to 26966
Data columns (total 11 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Unnamed: 0   26967 non-null  int64
1   carat        26967 non-null  float64
2   cut          26967 non-null  object
3   color        26967 non-null  object
4   clarity      26967 non-null  object
5   depth        26270 non-null  float64
6   table        26967 non-null  float64
7   x            26967 non-null  float64
8   y            26967 non-null  float64
9   z            26967 non-null  float64
10  price        26967 non-null  int64
dtypes: float64(6), int64(2), object(3)
memory usage: 2.3+ MB
```

Figure 1: Dataset Information

The figure below shows the description of the data. The range of each variable is the min value to the max value. Here it is observed that min values for variables x,y,z is zero which seems to be incorrect as each diamond will have an x,y,z as they are dimensions of the diamond. There seems to be a large difference in 75 percentile and max values which suggests there are outliers in the data. Carat seems to be one of the most important factors affecting the price of the diamond but a lot more analysis is required for that. The price of diamonds varies from 326 to 18818.

	Unnamed: 0	carat	depth	table	x	y	z	price
count	26967.000000	26967.000000	26270.000000	26967.000000	26967.000000	26967.000000	26967.000000	26967.000000
mean	13484.000000	0.798375	61.745147	57.456080	5.729854	5.733569	3.538057	3939.518115
std	7784.846891	0.477745	1.412860	2.232068	1.128516	1.168058	0.720624	4024.864666
min	1.000000	0.200000	50.800000	49.000000	0.000000	0.000000	0.000000	326.000000
25%	6742.500000	0.400000	61.000000	56.000000	4.710000	4.710000	2.900000	945.000000
50%	13484.000000	0.700000	61.800000	57.000000	5.690000	5.710000	3.520000	2375.000000
75%	20225.500000	1.050000	62.500000	59.000000	6.550000	6.540000	4.040000	5360.000000
max	26967.000000	4.500000	73.800000	79.000000	10.230000	58.900000	31.800000	18818.000000

Figure 2: Dataset Description

There are 697 null values present in the depth column

```

Unnamed: 0      0
carat           0
cut            0
color          0
clarity        0
depth         697
table          0
x              0
y              0
z              0
price          0
dtype: int64

```

Figure 3: Null Values

The Unnamed: 0 has been dropped as it is of no importance.

	carat	cut	color	clarity	depth	table	x	y	z	price
26	0.34	Ideal	D	SI1	NaN	57.0	4.50	4.44	2.74	803
86	0.74	Ideal	E	SI2	NaN	59.0	5.92	5.97	3.52	2501
117	1.00	Premium	F	SI1	NaN	59.0	6.40	6.36	4.00	5292
148	1.11	Premium	E	SI2	NaN	61.0	6.66	6.61	4.09	4177
163	1.00	Very Good	F	VS2	NaN	55.0	6.39	6.44	3.99	6340
...	...	...	...	...	...	...	...	...	...	...
26848	1.22	Very Good	H	VS1	NaN	59.0	6.91	6.85	4.29	7673
26854	1.29	Premium	I	VS2	NaN	58.0	7.12	7.03	4.27	6321
26879	0.51	Very Good	E	SI1	NaN	58.0	5.10	5.13	3.12	1343
26923	0.51	Ideal	D	VS2	NaN	57.0	5.12	5.09	3.18	1882
26960	1.10	Very Good	D	SI2	NaN	63.0	6.76	6.69	3.94	4361

697 rows × 10 columns

Figure 4: Records with Null Values

There are zeroes present in the dataset in variables 'x','y','z' which makes no sense as each diamond will have certain length, breadth and width so these records are not correct and the corresponding mean values of these variables have been imputed to the zeroes.

	carat	cut	color	clarity	depth	table	x	y	z	price
5821	0.71	Good	F	SI2	64.1	60.0	0.00	0.00	0.0	2130
6034	2.02	Premium	H	VS2	62.7	53.0	8.02	7.95	0.0	18207
10827	2.20	Premium	H	SI1	61.2	59.0	8.42	8.37	0.0	17265
12458	2.18	Premium	H	SI2	59.4	61.0	8.49	8.45	0.0	12631
12689	1.10	Premium	G	SI2	63.0	59.0	6.50	6.47	0.0	3696
17506	1.14	Fair	G	VS1	57.5	67.0	0.00	0.00	0.0	6381
18194	1.01	Premium	H	I1	58.1	59.0	6.66	6.60	0.0	3167
23758	1.12	Premium	G	I1	60.4	59.0	6.71	6.67	0.0	2383

Figure 5: Zero Value Records



There are 34 duplicate records in the dataset. The below figure shows the 34 duplicated records.

	carat	cut	color	clarity	depth	table	x	y	z	price
4756	0.35	Premium	J	VS1	62.4	58.0	5.67	5.64	3.53	949
6215	0.71	Good	F	SI2	64.1	60.0	0.00	0.00	0.00	2130
8144	0.33	Ideal	G	VS1	62.1	55.0	4.46	4.43	2.76	854
8919	1.52	Good	E	I1	57.3	58.0	7.53	7.42	4.28	3105
9818	0.35	Ideal	F	VS2	61.4	54.0	4.58	4.54	2.80	906
10473	0.79	Ideal	G	SI1	62.3	57.0	5.90	5.85	3.66	2898
10500	1.00	Premium	F	VVS2	60.6	54.0	6.56	6.52	3.96	8924
12894	1.21	Premium	D	SI2	62.5	57.0	6.79	6.71	4.22	6505
13547	0.43	Ideal	G	VS1	61.9	55.0	4.84	4.86	3.00	943
13783	0.79	Ideal	G	SI1	62.3	57.0	5.90	5.85	3.66	2898
14389	0.60	Premium	D	SI2	62.0	57.0	5.43	5.35	3.34	1196
14410	1.00	Very Good	D	SI1	63.1	56.0	6.34	6.30	3.99	5645
15798	0.90	Very Good	I	VS2	58.4	62.0	6.29	6.35	3.69	3334
16852	0.79	Ideal	G	SI1	62.3	57.0	5.90	5.85	3.66	2898
17263	1.04	Premium	I	SI2	62.0	57.0	6.53	6.47	4.03	3774
18025	1.51	Good	I	SI1	63.8	57.0	7.21	7.18	4.59	6046
18777	0.32	Premium	H	VS2	60.6	58.0	4.47	4.44	2.70	648
18837	1.01	Premium	H	VS1	61.2	61.0	6.44	6.41	3.93	5294
19731	0.30	Good	J	VS1	63.4	57.0	4.23	4.26	2.69	394
19877	2.01	Premium	I	VS2	60.3	62.0	8.13	8.06	4.89	15939
20301	0.30	Ideal	H	SI1	62.2	57.0	4.26	4.29	2.66	450
20760	1.80	Ideal	H	VS1	62.3	56.0	7.79	7.76	4.84	15105
22322	2.05	Premium	I	SI2	62.0	58.0	8.13	8.08	5.02	9850
22488	2.42	Premium	J	VS2	61.3	59.0	8.61	8.58	5.27	17168
22583	0.33	Ideal	F	IF	61.2	56.0	4.47	4.49	2.74	1240
23458	2.66	Good	H	SI2	63.8	57.0	8.71	8.65	5.54	16239
23564	1.50	Premium	F	SI2	58.5	60.0	7.52	7.48	4.39	7644
24351	2.50	Fair	H	SI2	64.9	58.0	8.46	8.43	5.48	13278
24816	1.50	Good	G	SI2	57.5	63.0	7.53	7.49	4.32	6006
25268	1.20	Premium	I	VS2	62.6	58.0	6.77	6.72	4.22	5699
25759	0.30	Ideal	G	IF	62.1	55.0	4.32	4.35	2.69	863
25941	0.51	Premium	F	SI2	58.1	59.0	5.26	5.24	3.05	1052
26191	2.54	Very Good	H	SI2	63.5	58.0	8.68	8.65	5.50	16353
26530	0.41	Ideal	G	IF	61.7	56.0	4.77	4.80	2.95	1367

Figure 6: Duplicated Records

## Univariate Analysis:

### Depth




Figure 7: Variable 'depth'

Observation: The boxplot of depth variable shows there are outliers on both sides. The distribution of the depth variable is almost normally distributed.

### Carat




Figure 8: Variable 'carat'

Observation: The boxplot of the carat variable shows a lot of outliers. The distribution plot shows that the distribution is right skewed.

### Table




Figure 9: Variable 'table'

Observation: The boxplot of the variable table shows that there are outliers on both ends. The distribution shows that table is right skewed.

X




Figure 10: Variable 'x'

Observation: The boxplot of the x variable shows that there are some outliers. The distribution plot shows that the distribution is right skewed.

Y




Figure 11: Variable 'y'

Observation: The boxplot of the y variable shows that there are very few outliers. The distribution plot shows that the distribution is right skewed.

Z




Figure 12: Variable 'Z'

Observation: The boxplot of the y variable shows that there are very few outliers. The distribution plot shows that the distribution is right skewed.

Price




Figure 13: Variable 'Price'

Observation: The boxplot of the price variable shows that there are some outliers. The distribution plot shows that the distribution is right skewed.

Bivariate Analysis




Figure 14: Boxplot

Boxplot between clarity and price shows that median for SI1, SI2, I1 is higher than the rest of them.




Figure 15: Boxplot

Boxplot between cut and price show that ideal cut has the least median value which is expected and premium has highest median which suggests premium cut will have high prices and ideal cut will low prices but the price ranges are seen spread out because there are different factors also affecting price.




Figure 16: Boxplot

Boxplot between color and price shows that H, I, J have higher median price than other colors.

## Multivariate Analysis

### Heatmap




Figure 17: Heatmap

The heatmap shows that there is a strong correlation between these variables:

- Carat and x
- Carat and y
- Carat and z

- Price and carat
- X and y
- Y and z
- X and z

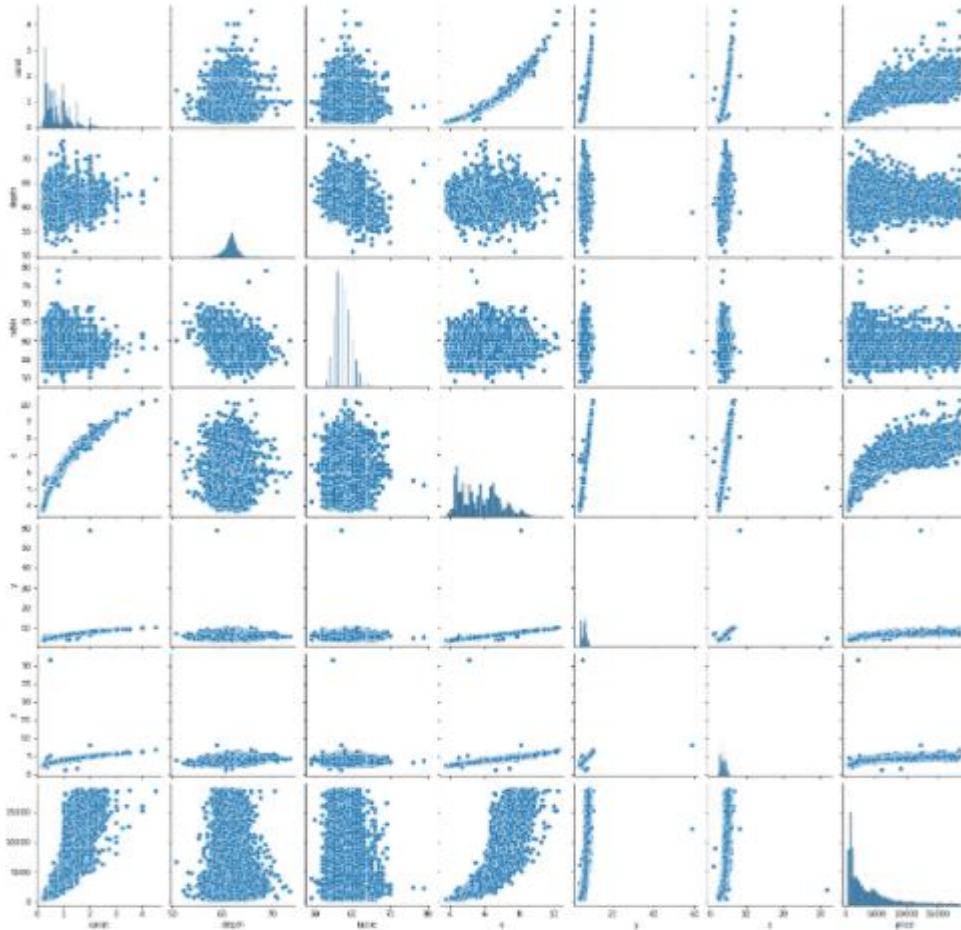


Figure 18: Pairplot

The pairplot also suggests the same as heatmap that there is good linear relation in carat and x,y,z.

To apply linear regression, the datatypes of variables have to be int or float. So, to convert them into numerical values one hot encoding is used and the variables have been increased and all have been converted to numeric variables.

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 26933 entries, 0 to 26966
Data columns (total 24 columns):
#   Column                Non-Null Count  Dtype
---  -
0   carat                 26933 non-null  float64
1   depth                 26236 non-null  float64
2   table                 26933 non-null  float64
3   x                     26931 non-null  float64
4   y                     26931 non-null  float64
5   z                     26925 non-null  float64
6   price                 26933 non-null  int64
7   cut_Good              26933 non-null  uint8
8   cut_Ideal             26933 non-null  uint8
9   cut_Premium          26933 non-null  uint8
10  cut_Very Good         26933 non-null  uint8
11  clarity_IF            26933 non-null  uint8
12  clarity_SI1           26933 non-null  uint8
13  clarity_SI2           26933 non-null  uint8
14  clarity_VS1           26933 non-null  uint8
15  clarity_VS2           26933 non-null  uint8
16  clarity_VVS1          26933 non-null  uint8
17  clarity_VVS2          26933 non-null  uint8
18  color_E               26933 non-null  uint8
19  color_F               26933 non-null  uint8
20  color_G               26933 non-null  uint8
21  color_H               26933 non-null  uint8
22  color_I               26933 non-null  uint8
23  color_J               26933 non-null  uint8
dtypes: float64(6), int64(1), uint8(17)
memory usage: 3.1 MB

```

Figure 19: Columns after Label Encoding

Linear regression is very sensitive to outliers and all these outliers have been treated to the upper and lower whisker of the boxplot. The figure below shows that there are no outliers left in the dataset. However, the linear regression has been applied on the dataset with both the outliers present and without the outliers.

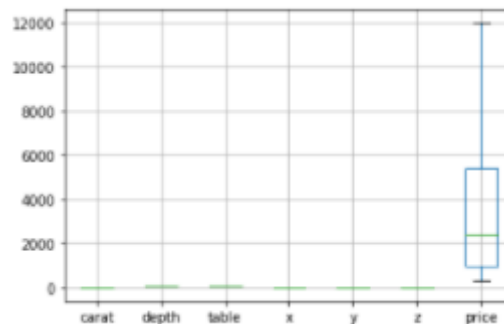


Figure 20: Boxplot After Outlier Treatment

## Question 1.2

Impute null values if present, also check for the values which are equal to zero. Do they have any meaning or do we need to change them or drop them? Check for the possibility of combining the sub levels of an ordinal variables and take actions accordingly. Explain why you are combining these sub levels with appropriate reasoning.

## Answer 1.2

There are 697 null values present in depth variable.

```
Unnamed: 0      0
carat           0
cut             0
color           0
clarity         0
depth          697
table           0
x              0
y              0
z              0
price           0
dtype: int64
```

Figure 21: Null Values

There are some values which are equal to zero in x,y,z variables. These have been replaced by null values as these values cannot be zero as each diamond would have a length, breadth and height.

```
carat           0
depth          697
table           0
x               2
y               2
z               8
price           0
cut_Good        0
cut_Ideal       0
cut_Premium     0
cut_Very Good   0
clarity_IF      0
clarity_SI1     0
clarity_SI2     0
clarity_VS1     0
clarity_VS2     0
clarity_VVS1    0
clarity_VVS2    0
color_E         0
color_F         0
color_G         0
color_H         0
color_I         0
color_J         0
dtype: int64
```

Figure 22: Null Values after Data Modification



The null values in depth, x, y, z have been imputed by the mean values of these variables. The below figure shows that there are no null values present in the dataset

```
carat      0
depth      0
table      0
x          0
y          0
z          0
price      0
cut_Good   0
cut_Ideal  0
cut_Premium 0
cut_Very Good 0
clarity_IF 0
clarity_SI1 0
clarity_SI2 0
clarity_VS1 0
clarity_VS2 0
clarity_VVS1 0
clarity_VVS2 0
color_E    0
color_F    0
color_G    0
color_H    0
color_I    0
color_J    0
dtype: int64
```

Figure 23: Null values After Imputation

Some models have been built in order to increase the performance by combining certain sub levels of variables based on the domain knowledge about diamonds.

- Color D, E, F have been combined to a Colorless level.
- Color G,H,I,J have been combined to a Near Colorless level.
- Clarity I1, SI1, SI2 have been combined to an Impure level.
- Clarity VS1, VS2 have been combined to a Slightly-Impure level.
- Clarity VVS1, VVS2 have been combined to a Very\_Slightly\_Impure level.
- Clarity IF have been combined to a No\_Impurities level.

Same technique of getting dummy variables using one hot encoding has been used to convert categorical variables to numeric variables.

### Question 1.3

Encode the data (having string values) for Modelling. Split the data into train and test (70:30). Apply Linear regression using scikit learn. Perform checks for significant variables using appropriate method from statsmodel. Create multiple models and check the performance of Predictions on Train and Test sets using Rsquare, RMSE & Adj Rsquare. Compare these models and select the best one with appropriate reasoning.

### Answer 1.3

The data having string values have been encoded using the one hot encoding method. This technique converts the categorical variables into numerical variables by making new columns equal to the levels in each original column and dropping the first column. The below figure shows the new columns that have been created and all the columns have been converted into numerical data types.

cut_Good	cut_Ideal	cut_Premium	...	clarity_VS1	clarity_VS2	clarity_VVS1	clarity_VVS2	color_E	color_F	color_G	color_H	color_I	color_J
0	1	0	...	0	0	0	0	1	0	0	0	0	0
0	0	1	...	0	0	0	0	0	0	1	0	0	0
0	0	0	...	0	0	0	1	1	0	0	0	0	0
0	1	0	...	1	0	0	0	0	1	0	0	0	0
0	1	0	...	0	0	1	0	0	1	0	0	0	0

Figure 24: One Hot Encoding

The data has been splitted into 70% train and 30% test data. Linear Regression analysis is applied on train data using scikit learn library. For further analysis Linear regression using statsmodel has been applied to get adjusted R squared and other parameters to decide the best model that can be used for the current dataset.

### Model 1

The first model is built using all the columns and without dropping any of them.

Using scikit learn library applying Linear Regression Model.

Model score for a regression model is nothing but r squared.

- Score for Training data is 0.940
- Score for Test data is 0.942
- Root mean squared error for training data is 846.94
- Root mean squared error for test data is 836.75

This represents the best fit line. The model is right fit and train and test data are in sync

Using stats model applying Linear Regression model as it gives more widened performance metrics. R squared and adjusted r squared value is 0.940 which is very good.

```

=====
Dep. Variable:          price      R-squared:          0.940
Model:                  OLS        Adj. R-squared:      0.940
Method:                 Least Squares  F-statistic:        1.289e+04
Date:                   Wed, 12 Jan 2022  Prob (F-statistic):    0.00
Time:                   16:33:22      Log-Likelihood:     -1.5385e+05
No. Observations:       18853        AIC:                3.078e+05
Df Residuals:           18829        BIC:                3.079e+05
Df Model:               23
Covariance Type:        nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
Intercept	-3387.1169	748.747	-4.417	0.000	-4774.728	-1839.506
carat	9177.8162	77.380	118.607	0.000	9026.145	9329.488
depth	10.7415	10.394	1.033	0.301	-9.632	31.115
table	-18.9513	3.842	-4.933	0.000	-26.482	-11.421
x	-1086.8989	123.094	-8.830	0.000	-1328.174	-845.624
y	967.9215	124.393	7.781	0.000	724.101	1211.742
z	-577.3369	129.075	-4.473	0.000	-830.336	-324.337
cut_Good	460.9612	44.366	10.390	0.000	374.000	547.922
cut_Ideal	698.0794	43.141	16.181	0.000	613.520	782.639
cut_Premium	659.6221	41.422	15.925	0.000	578.432	740.812
cut_Very_Good	590.4956	42.407	13.924	0.000	507.373	673.618
clarity_IF	3992.6455	66.195	60.316	0.000	3862.897	4122.394
clarity_SI1	2510.2269	56.650	44.311	0.000	2399.188	2621.266
clarity_SI2	1674.8881	56.927	29.422	0.000	1563.305	1786.471
clarity_VS1	3333.5315	57.765	57.709	0.000	3220.308	3446.755
clarity_VS2	3029.8539	56.976	53.178	0.000	2918.176	3141.532
clarity_VVS1	3761.3204	61.011	61.650	0.000	3641.734	3880.907
clarity_VVS2	3746.8059	59.432	63.043	0.000	3630.314	3863.298
color_E	-181.2250	22.737	-7.971	0.000	-225.791	-136.659
color_F	-256.2452	23.206	-11.042	0.000	-301.730	-210.760
color_G	-428.1865	22.515	-19.018	0.000	-472.318	-384.055
color_H	-856.3686	24.094	-35.543	0.000	-903.594	-809.143
color_I	-1325.3521	26.791	-49.469	0.000	-1377.865	-1272.839
color_J	-1929.0572	33.026	-58.410	0.000	-1993.792	-1864.322

```

=====
Omnibus:                 4775.536   Durbin-Watson:          1.983
Prob(Omnibus):           0.000     Jarque-Bera (JB):       17747.449
Skew:                    1.234     Prob(JB):               0.00
Kurtosis:                 7.062     Cond. No.               1.04e+04
=====
Notes:
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
[2] The condition number is large, 1.04e+04. This might indicate that there are
strong multicollinearity or other numerical problems.

```

Figure 25: Model 1 Regression Analysis

Hypothesis testing for Linear Regression – The null hypothesis states that there is no relation between the dependent variable Price and other independent variables. Looking at the summary table above, all the P values are less than 0.05 or at 95% confidence level we can say that the variables have a direct impact on the price variable except the depth variable which has p value greater than 0.05. For depth null hypothesis cannot be rejected and it is inferred that depth has no effect on depth variable. Carat and clarity variables seem to impact the price rise positively, surprisingly; color of the stones is reducing the price increase. We can study our model further to see if we can reduce multicollinearity, if present to get the correct coefficients.

The below scatter plot is between actual price on x axis and predicted price on y axis. The scatter plot shows a strong linear relationship between them meaning the model gives out a really good prediction.




Figure 26: Scatter Plot Between Actual and Predicted Price

The below image is variance inflation factor which gives which factor is the reason for multicollinearity in the data. Here it is observed that depth, table, x, y, z are the most contributing factors to multicollinearity and hence further models have to be made such to address this multicollinearity. Multicollinearity affects the coefficients of independent variables and they may not be correct in predicting the target variable so this multicollinearity has to be reduced to get the correct coefficients.

```
carat ---> 122.83209929169017
depth ---> 1348.5747037839517
table ---> 978.7910961905317
x ---> 11960.44592119137
y ---> 11486.794645087222
z ---> 3179.596094744831
cut_Good ---> 4.4914016280632785
cut_Ideal ---> 18.016728432469563
cut_Premium ---> 10.83115735472792
cut_Very Good ---> 10.01637987063325
clarity_IF ---> 3.655387429980751
clarity_SI1 ---> 19.688570072296287
clarity_SI2 ---> 13.82199801875791
clarity_VS1 ---> 12.746420794843807
clarity_VS2 ---> 18.44735046547165
clarity_VVS1 ---> 6.431913987513056
clarity_VVS2 ---> 8.379930444886524
color_E ---> 2.480829824573142
color_F ---> 2.448039716583097
color_G ---> 2.796050687532723
color_H ---> 2.305043749504993
color_I ---> 1.9312909945309655
color_J ---> 1.5142333068255323
```

Figure 27: Variance Inflation Factor

## Model 2:

The second model has been built by dropping the Depth Variable as it had a p value greater than 0.05.

First using scikitlearn library Linear Regression Analysis is applied

- Score for Training data is 0.940
- Score for Test data is 0.942
- Root mean squared error for training data is 846.96
- Root mean squared error for test data is 836.81

There seems to no change in the performance metrics from the first model.

Using statsmodel to build the Linear Regression model. Although there is no change in r squared and adjusted r squared the f statistic has improved considerably.

OLS Regression Results						
Dep. Variable:	price	R-squared:	0.940			
Model:	OLS	Adj. R-squared:	0.940			
Method:	Least Squares	F-statistic:	1.347e+04			
Date:	Wed, 12 Jan 2022	Prob (F-statistic):	0.00			
Time:	19:11:54	Log-Likelihood:	-1.5385e+05			
No. Observations:	18853	AIC:	3.077e+05			
Df Residuals:	18830	BIC:	3.079e+05			
Df Model:	22					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
Intercept	-2583.1184	264.237	-9.776	0.000	-3101.046	-2065.191
carat	9186.0017	76.974	119.340	0.000	9035.126	9336.877
table	-19.7726	3.759	-5.260	0.000	-27.140	-12.405
x	-1112.1452	120.646	-9.218	0.000	-1348.622	-875.669
y	923.6574	116.786	7.909	0.000	694.746	1152.569
z	-470.0193	76.660	-6.131	0.000	-620.279	-319.760
cut_Good	463.4877	44.299	10.463	0.000	376.659	550.317
cut_Ideal	694.1161	42.970	16.154	0.000	609.891	778.341
cut_Premium	655.7942	41.256	15.896	0.000	574.929	736.659
cut_Very_Good	588.7693	42.375	13.894	0.000	505.711	671.827
clarity_IF	3992.6306	66.195	60.316	0.000	3862.882	4122.380
clarity_SI1	2511.9543	56.625	44.361	0.000	2400.963	2622.945
clarity_SI2	1676.2228	56.913	29.452	0.000	1564.669	1787.777
clarity_VS1	3334.3531	57.759	57.729	0.000	3221.140	3447.566
clarity_VS2	3031.1913	56.961	53.215	0.000	2919.542	3142.841
clarity_VVS1	3761.8847	61.009	61.662	0.000	3642.302	3881.467
clarity_VVS2	3747.7677	59.425	63.067	0.000	3631.289	3864.246
color_E	-181.4035	22.736	-7.979	0.000	-225.968	-136.839
color_F	-256.2485	23.206	-11.043	0.000	-301.734	-210.763
color_G	-427.9656	22.514	-19.009	0.000	-472.096	-383.836
color_H	-855.8721	24.089	-35.530	0.000	-903.089	-808.656
color_I	-1324.6763	26.783	-49.459	0.000	-1377.174	-1272.178
color_J	-1928.6259	33.024	-58.401	0.000	-1993.356	-1863.896
Omnibus:	4778.989	Durbin-Watson:	1.983			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	17785.860			
Skew:	1.234	Prob(JB):	0.00			
Kurtosis:	7.068	Cond. No.	2.57e+03			

Figure 28: Model 2 Regression Analysis

The scatter plot between actual and predicted price also shows a linear graph which tells model is a good model for predicting price.




Figure 29: Scatter Plot Between Actual and Predicted Price

The Variance inflation factor below shows high value for x, y, z variable which means they are the contributing factors to multicollinearity in the data. X, y, z variables have a very strong relation with carat variable as carat variable increases it is bound that x, y, z will increase as they are the size measures.

```
carat ---> 102.22998573256453
table ---> 343.67178696032374
x ---> 11907.66520065791
y ---> 11068.192853143442
z ---> 1592.6925974471997
cut_Good ---> 4.364430991175007
cut_Ideal ---> 16.032354095929502
cut_Premium ---> 10.552785297072413
cut_Very Good ---> 9.581139499530757
clarity_IF ---> 3.507794731999077
clarity_SI1 ---> 18.901145707067478
clarity_SI2 ---> 13.307124818568566
clarity_VS1 ---> 12.220903273004456
clarity_VS2 ---> 17.66874831948798
clarity_VVS1 ---> 6.149324444822197
clarity_VVS2 ---> 8.015137230464944
color_E ---> 2.478314015891665
color_F ---> 2.4444749919155604
color_G ---> 2.7914199796781376
color_H ---> 2.296962880031655
color_I ---> 1.92588279414264
color_J ---> 1.5119985802546922
```

Figure 30: Variance Inflation Factor

### Model 3

The next model is built to reduce the multicollinearity in the data by further dropping variables x, y, z

First using scikitlearn library Linear Regression Analysis is applied

- Score for Training data is 0.939
- Score for Test data is 0.940
- Root mean squared error for training data is 853.72
- Root mean squared error for test data is 844.189

RMSE has increased by a little in this model but there is not any significant difference and train and test data are in line with each other.

Using statsmodel to build the Linear Regression model. Although there is no change in r squared and adjusted r squared the f statistic has improved considerably. All the p values are below 0 which shows that all the independent variables in this model have a significant impact on target variable price.

```

=====
Dep. Variable:          price    R-squared:          0.939
Model:                  OLS      Adj. R-squared:       0.939
Method:                 Least Squares    F-statistic:       1.534e+04
Date:                   Wed, 12 Jan 2022    Prob (F-statistic): 0.00
Time:                   19:12:02      Log-Likelihood:    -1.5400e+05
No. Observations:       18853      AIC:              3.080e+05
Df Residuals:           18833      BIC:              3.082e+05
Df Model:                19
Covariance Type:        nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
Intercept	-4717.6694	217.998	-21.641	0.000	-5144.965	-4290.373
carat	8039.2039	15.515	518.159	0.000	8008.793	8069.615
table	-16.9206	3.550	-4.766	0.000	-23.880	-9.961
cut_Good	551.9846	43.215	12.773	0.000	467.280	636.689
cut_Ideal	788.4697	40.630	19.406	0.000	708.831	868.109
cut_Premium	717.1073	39.820	18.009	0.000	639.057	795.157
cut_Very_Good	700.3553	40.311	17.374	0.000	621.343	779.368
clarity_IF	4103.3041	66.308	61.882	0.000	3973.334	4233.274
clarity_SI1	2549.2986	56.950	44.763	0.000	2437.671	2660.927
clarity_SI2	1711.2926	57.250	29.892	0.000	1599.077	1823.508
clarity_VS1	3391.3872	58.031	58.441	0.000	3277.642	3505.133
clarity_VS2	3086.3556	57.247	53.913	0.000	2974.146	3198.565
clarity_VVS1	3864.0629	61.159	63.180	0.000	3744.185	3983.941
clarity_VVS2	3829.7494	59.635	64.220	0.000	3712.859	3946.639
color_E	-183.4199	22.915	-8.004	0.000	-228.335	-138.504
color_F	-266.7587	23.379	-11.410	0.000	-312.585	-220.933
color_G	-438.3437	22.680	-19.327	0.000	-482.799	-393.888
color_H	-855.5083	24.265	-35.257	0.000	-903.070	-807.947
color_I	-1312.9075	26.973	-48.676	0.000	-1365.776	-1260.039
color_J	-1913.2713	33.262	-57.520	0.000	-1978.469	-1848.074

```

=====
Omnibus:                4285.730    Durbin-Watson:       1.980
Prob(Omnibus):           0.000      Jarque-Bera (JB):    13193.004
Skew:                    1.168      Prob(JB):            0.00
Kurtosis:                6.367      Cond. No.             2.07e+03
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 2.07e+03. This might indicate that there are strong multicollinearity or other numerical problems.

Figure 31: Model 3 Regression Analysis

The scatter plot between actual and predicted price also shows a linear graph which tells model is a good model for predicting price.




Figure 32: Scatter Plot Between Actual and Predicted Price

The Variance Inflation factor shows only table is the factor remaining which is showing some multicollinearity. All other variables contribute very little to the multicollinearity.

```
carat ---> 5.223684588164773
table ---> 97.56394186616357
cut_Good ---> 4.1312866621866124
cut_Ideal ---> 14.335956924499984
cut_Premium ---> 9.885733103057277
cut_Very Good ---> 8.70333130825681
clarity_IF ---> 3.481310724007501
clarity_SI1 ---> 18.559797736467978
clarity_SI2 ---> 13.097823682775097
clarity_VS1 ---> 12.053613938763373
clarity_VS2 ---> 17.42144406697902
clarity_VVS1 ---> 6.101273452791654
clarity_VVS2 ---> 7.9337889552988
color_E ---> 2.4758946280948533
color_F ---> 2.4352303532503305
color_G ---> 2.7781168170164494
color_H ---> 2.290612795680401
color_I ---> 1.9225371559098865
color_J ---> 1.509672238133545
```

Figure 33: Variance Inflation Factor

#### Model 4

This model is built by dropping the table variable further to get rid of the multicollinearity.

First using scikitlearn library Linear Regression Analysis is applied

- Score for Training data is 0.939
- Score for Test data is 0.940
- Root mean squared error for training data is 854.23
- Root mean squared error for test data is 844.833

RMSE has increased by a little in this model but there is not any significant difference. The train and test data are in line.

Using statsmodel to build the Linear Regression model. Although there is no change in r squared and adjusted r squared the f statistic has improved considerably. All the p values are below 0 which shows that all the independent variables in this model have a significant impact on target variable price.



```

=====
Dep. Variable:          price      R-squared:          0.939
Model:                  OLS       Adj. R-squared:       0.939
Method:                 Least Squares   F-statistic:        1.618e+04
Date:                  Wed, 12 Jan 2022   Prob (F-statistic):  0.00
Time:                  19:12:03    Log-Likelihood:     -1.5401e+05
No. Observations:      18853       AIC:                3.081e+05
Df Residuals:          18834       BIC:                3.082e+05
Df Model:              18
Covariance Type:       nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
Intercept	-5707.2444	66.430	-85.914	0.000	-5837.453	-5577.035
carat	8032.1135	15.452	519.799	0.000	8001.826	8062.402
cut_Good	552.8581	43.239	12.786	0.000	468.105	637.611
cut_Ideal	833.8164	39.523	21.097	0.000	756.348	911.285
cut_Premium	717.4752	39.843	18.008	0.000	639.380	795.570
cut_Very_Good	713.1301	40.245	17.720	0.000	634.247	792.013
clarity_IF	4107.9141	66.339	61.923	0.000	3977.883	4237.945
clarity_SI1	2551.6955	56.981	44.781	0.000	2440.008	2663.383
clarity_SI2	1712.3615	57.283	29.893	0.000	1600.082	1824.641
clarity_VS1	3394.0323	58.062	58.456	0.000	3280.226	3507.838
clarity_VS2	3088.7704	57.278	53.926	0.000	2976.501	3201.040
clarity_VVS1	3865.5425	61.194	63.169	0.000	3745.597	3985.488
clarity_VVS2	3830.9495	59.669	64.203	0.000	3713.993	3947.906
color_E	-185.2647	22.925	-8.081	0.000	-230.200	-140.330
color_F	-266.4984	23.393	-11.392	0.000	-312.351	-220.646
color_G	-437.3437	22.692	-19.273	0.000	-481.823	-392.864
color_H	-853.8460	24.276	-35.172	0.000	-901.430	-806.262
color_I	-1312.3437	26.988	-48.627	0.000	-1365.242	-1259.445
color_J	-1913.8403	33.281	-57.505	0.000	-1979.075	-1848.606

```

=====
Omnibus:                4299.918      Durbin-Watson:       1.980
Prob(Omnibus):          0.000        Jarque-Bera (JB):    13189.906
Skew:                   1.173        Prob(JB):            0.00
Kurtosis:               6.359        Cond. No.            38.6
=====
Notes:
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```

Figure 34: Model 4 Regression Analysis

The problem of multicollinearity is removed as the VIF of this model indicates there is no multicollinearity in the data.

```

carat ---> 4.023963552948521
cut_Good ---> 3.5020910476346123
cut_Ideal ---> 12.444657871405733
cut_Premium ---> 8.172657501145094
cut_Very Good ---> 7.362893118877587
clarity_IF ---> 2.111573704227221
clarity_SI1 ---> 8.582797349957549
clarity_SI2 ---> 6.341835689528017
clarity_VS1 ---> 5.81765676176599
clarity_VS2 ---> 8.132009944538451
clarity_VVS1 ---> 3.2360618789461135
clarity_VVS2 ---> 3.994479451757277
color_E ---> 2.3711536081820666
color_F ---> 2.340035283675498
color_G ---> 2.6842016288544035
color_H ---> 2.221751381261127
color_I ---> 1.8871709780128412
color_J ---> 1.4941185157588375

```

Figure 35: Variance Inflation Factor

## Model 5:

This model is built using feature modeling and combining different ordinal sub levels into 1 to check if there is any improvement in the performance metrics. The combined levels have been discussed in the previous question.

```

carat      0
depth      0
table      0
x          0
y          0
z          0
price      0
cut_Good   0
cut_Ideal  0
cut_Premium 0
cut_Very Good 0
clarity_No_Impurities 0
clarity_Slightly_Impure 0
clarity_Very_Slightly_Impure 0
color_Near_Colorless 0
dtype: int64

```

Figure 36: Modified Columns

First using scikitlearn library Linear Regression Analysis is applied

- Score for Training data is 0.920
- Score for Test data is 0.920
- Root mean squared error for training data is 979.76
- Root mean squared error for test data is 979.80

RMSE has increased and Rsquared value has dropped which tells this is not a good model.

```

=====
Dep. Variable: price      R-squared: 0.920
Model: OLS              Adj. R-squared: 0.920
Method: Least Squares   F-statistic: 1.549e+04
Date: Sun, 16 Jan 2022   Prob (F-statistic): 0.00
Time: 21:57:35          Log-Likelihood: -1.5660e+05
No. Observations: 18853 AIC: 3.132e+05
Df Residuals: 18838     BIC: 3.133e+05
Df Model: 14
Covariance Type: nonrobust
=====
              coef      std err          t      P>|t|      [0.025      0.975]
-----
Intercept    -1825.5543    864.108     -2.113    0.035    -3519.284    -131.825
carat         8672.2225    88.950     97.495    0.000    8497.872    8846.573
depth         18.2815     12.007      1.523    0.128     -5.253     41.816
table        -25.6301      4.441     -5.772    0.000    -34.334    -16.926
x            -1178.1551    142.155     -8.288    0.000   -1456.792    -899.518
y            1245.1260    143.490      8.677    0.000     963.873    1526.379
z            -731.4832    149.058     -4.907    0.000   -1023.651    -439.316
cut_Good       698.6566     51.005     13.698    0.000     598.683     798.631
cut_Ideal     951.3730     49.527     19.209    0.000     854.295    1048.451
cut_Premium   927.1785     47.499     19.520    0.000     834.077    1020.280
cut_Very_Good 845.2556     48.655     17.372    0.000     749.887     940.624
clarity_No_Impurities 1873.9209     42.471     44.123    0.000    1790.674    1957.168
clarity_Slightly_Impure 1003.5859     16.470     60.935    0.000     971.304    1035.868
clarity_Very_Slightly_Impure 1615.4711     22.269     72.543    0.000    1571.821    1659.121
color_Near_Colorless -660.8148     14.838    -44.534    0.000    -689.899    -631.730
=====
Omnibus: 3020.949   Durbin-Watson: 2.001
Prob(Omnibus): 0.000   Jarque-Bera (JB): 12340.520
Skew: 0.750   Prob(JB): 0.00
Kurtosis: 6.669   Cond. No. 1.04e+04
=====

Notes:
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
[2] The condition number is large, 1.04e+04. This might indicate that there are
strong multicollinearity or other numerical problems.

```

Figure 37: Model 5 Regression Analysis

Using statsmodel to build the Linear Regression model. R squared and adjusted r squared value has decreased. Depth has p value greater than alpha so there is not enough evidence to reject null hypothesis.

This model is not a good model as compared to the previous models.

For deciding best model in linear regression, r squared value should be high and RMSE should be low as possible. Further the adjusted r squared and other metrics are checked using statsmodel. To get the correct coefficients multicollinearity should be reduced.

Model Comparison:

	Model1 Train	Model1 Test	Model2 Train	Model2 Test	Model3 Train	Model3 Test	Model4 Train	Model4 Test	Model5 Train	Model5 Test
R Squared	0.940	0.941	0.940	0.941	0.939	0.940	0.939	0.940	0.920	0.920
RMSE	846.94	836.75	846.96	836.81	853.72	844.18	854.23	844.23	979.76	979.80
Adjusted R squared	0.94		0.94		0.939		0.939		0.92	
Multicollinearity	Yes		Yes		Yes		No		Yes	

Table 1: Model comparison

Model 4 has very good r squared and adjusted r squared and it has the highest f statistic. The coefficients can be well explained for this model as there is no multicollinearity present for this model.

## Question 1.4

Inference: Basis on these predictions, what are the business insights and recommendations.

## Answer 1.4

The 5 most Important factors affecting the price of diamonds are:

Carat, clarity\_IF, clarity\_VVS1, clarity\_VVS2, clarity\_VS1 are the 5 attributes that are most important for this dataset.

The final regression equation is:

```
(-5787.24) * Intercept + (8032.11) * carat + (552.86) * cut_Good + (833.82) * cut_Ideal + (717.48) * cut_Premium + (713.13) * cut_Very_Good + (4187.91) * clarity_IF + (2551.7) * clarity_SI1 + (1712.36) * clarity_SI2 + (3394.83) * clarity_VS1 + (3888.77) * clarity_VS2 + (3865.54) * clarity_VVS1 + (3838.95) * clarity_VVS2 + (-185.26) * color_E + (-266.5) * color_F + (-437.34) * color_G + (-853.85) * color_H + (-1312.34) * color_I + (-1913.84) * color_J +
```

Figure 38: Linear Regression Equation

When carat increases by 1 unit, price increases by 8032.11 units, keeping all other predictors constant. Clarity is a strong predictor as clarity\_IF, clarity\_SI1, clarity\_SI2, clarity\_VS1, clarity\_VS2, clarity\_VVS1, clarity\_VVS2 have high coefficients.

There are also some negative co-efficient values, for instance, color has all its coefficients negative. The colorless colors like D, E, F contribute very less negatively comparatively to colors G, H, I, J.

Cut which are of premium and very good quality seem to have a higher positive impact. Although Ideal cut also has a very high coefficient.

Recommendations

- The colors H, I, J of the diamonds should be avoided as they contribute very negatively to the price so colors D, E, F, G should be used.
- The company should focus on clarity and carat of the diamonds so as to increase the price of the diamond.
- Company can introduce a lottery system for people who buy expensive diamonds and give them a chance to receive some gifts.
- The customers will be of different financial backgrounds so all type of diamonds from low range to high range should be available in stock. A further analysis can be done by the company based on the customer's financial backgrounds to understand which diamonds they are likely to buy.

## Problem 2: Logistic Regression and LDA

You are hired by a tour and travel agency which deals in selling holiday packages. You are provided details of 872 employees of a company. Among these employees, some opted for the package and some didn't. You have to help the company in predicting whether an employee will opt for the package or not on the basis of the information given in the data set. Also, find out the important factors on the basis of which the company will focus on particular employees to sell their packages.

Variable Name	Description
Holiday_Package	Opted for Holiday Package yes/no?
Salary	Employee salary
age	Age in years
edu	Years of formal education
no_young_children	The number of young children (younger than 7 years)
no_older_children	Number of older children
foreign	foreigner Yes/No

### Question 2.1

Data Ingestion: Read the dataset. Do the descriptive statistics and do null value condition check, write an inference on it? Perform Univariate and Bivariate Analysis. Do exploratory data analysis.

### Answer 2.1

The dataset consists of 872 rows and 8 columns. 2 columns are of datatype object namely Holliday\_Package and foreign. All other columns are of continuous nature and datatype as int or float.

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 872 entries, 0 to 871
Data columns (total 8 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   Unnamed: 0            872 non-null   int64
1   Holliday_Package      872 non-null   object
2   Salary                872 non-null   int64
3   age                  872 non-null   int64
4   educ                 872 non-null   int64
5   no_young_children    872 non-null   int64
6   no_older_children    872 non-null   int64
7   foreign               872 non-null   object
dtypes: int64(6), object(2)
memory usage: 54.6+ KB
```

Figure 39: Dataset Information

The figure below shows the description of the data. The range of each variable is the min value to the max value. There seems to be a large difference in 75 percentile and max values for Salary,

no\_young\_children, no\_older\_children variable which suggests there are outliers in the data. Number of older children varies from 0 to 6 and number of younger children vary from 0 to 3.

	Unnamed: 0	Salary	age	educ	no_young_children	no_older_children
count	872.000000	872.000000	872.000000	872.000000	872.000000	872.000000
mean	436.500000	47729.172018	39.955275	9.307339	0.311927	0.982798
std	251.889014	23418.688531	10.551675	3.036259	0.612870	1.088788
min	1.000000	1322.000000	20.000000	1.000000	0.000000	0.000000
25%	218.750000	35324.000000	32.000000	8.000000	0.000000	0.000000
50%	436.500000	41903.500000	39.000000	9.000000	0.000000	1.000000
75%	654.250000	53489.500000	48.000000	12.000000	0.000000	2.000000
max	872.000000	238981.000000	62.000000	21.000000	3.000000	6.000000

Figure 40: Dataset Description

There are no null values present in the dataset, therefore no imputation is required.

```

Unnamed: 0      0
Holiday_Package  0
Salary          0
age            0
educ           0
no_young_children  0
no_older_children  0
foreign        0
dtype: int64

```

Figure 41: Null Values

There are no duplicates in the data.

The Unnamed: 0 column is of no use and therefore it has been dropped from the dataset.

The percentage of employees who have opted for the package is 45%. The data is a balanced between people who have opted and not opted for the package.

```

no      0.540138
yes     0.459862
Name: Holiday_Package, dtype: float64

```

Figure 42: Holiday\_Package Value Count

There are 216 employees out of the total records who are foreigners.

```

no      656
yes     216
Name: foreign, dtype: int64

```

Figure 43: Foreign Value Count

## Univariate Analysis

## Salary




Figure 44: Variable 'Salary'

The boxplot for salary variable shows that there are a lot of outliers. The distribution plot shows that salary variable is right skewed.

## Age




Figure 45: Variable 'Age'

The boxplot for Age variable shows that there are no outliers. The distribution plot shows that age variable is almost normally distributed.

## Educ




Figure 46: Variable 'Educ'

The boxplot of Educ variable shows that there are outliers on both ends. The distribution plot shows that it is left skewed. Educ is a discrete numeric variable

No\_young\_children




Figure 47: Variable 'No\_young\_children'

The boxplot for no\_young\_children variable shows that there are a lot of outliers(anything other than value 0 is considered an outlier for this variable) The distribution plot shows that no\_young\_children variable is right skewed. This no\_young\_children is a discrete numeric variable.

No\_older\_children




Figure 48; Variable 'No\_older\_children'

The boxplot for no\_older\_children variable shows that there are a few outliers. The distribution plot shows that no\_older\_children variable is right skewed. This no\_older\_children is a discrete numeric variable.

Bivariate Analysis



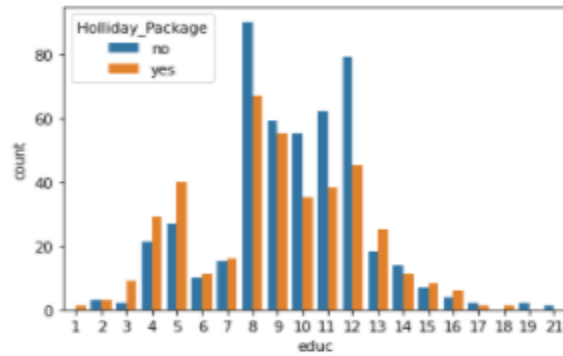


Figure 49: Count plot of Educ

The countplot for educ variable with a hue of Holliday\_Package variable shows that it widely distributed and most of the employees have educ between 4 to 13 years and ratio between opted and not opted is almost 60/40.

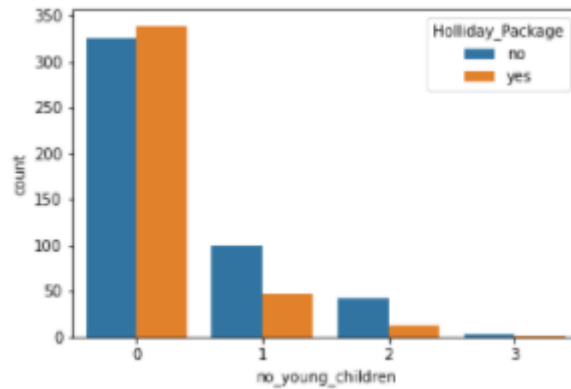


Figure 50: Count Plot of no\_young\_children

The countplot for no\_young\_children shows that people with 0 children have chosen the package more and people with children have hardly chosen the package which is expected as people do not tend to travel with young children.

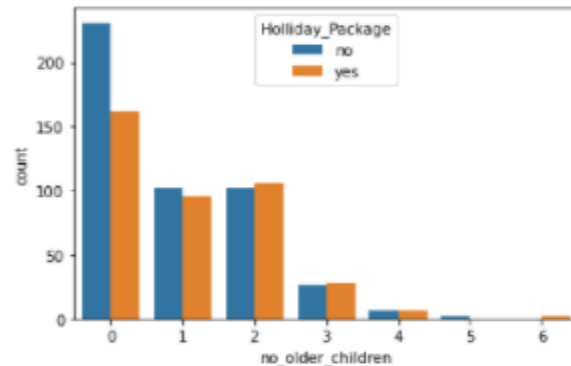


Figure 51: Count plot of no\_older\_children

The countplot for no\_older\_children with hue of Holliday\_Package shows how many employees have chosen and not chosen the package. From the plot it seems that proportion that most employees have a smaller number of older children and the proportion between chosen/not chosen improves as the number of children increases which can be due to the reason that older children being independent and the employees can go for holiday.

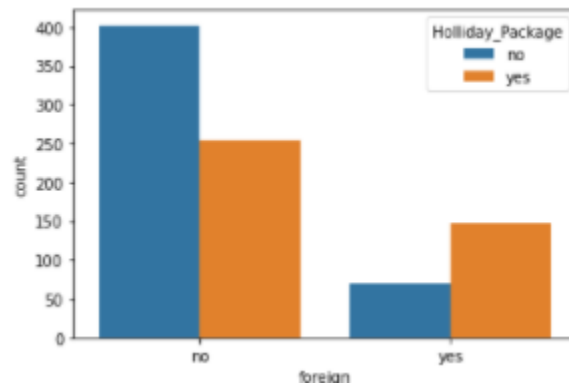


Figure 52: Count Plot of foreign

The counplot for foreign with hue of Holliday\_Package shows that employees who are foreigners have opted for the holiday package more that the employees who are not foreigners.




Figure 53: Heatmap

The heatmap shows there is not any high correlation in any of the variables. However, there is some correlation between no\_young\_children and age variable.

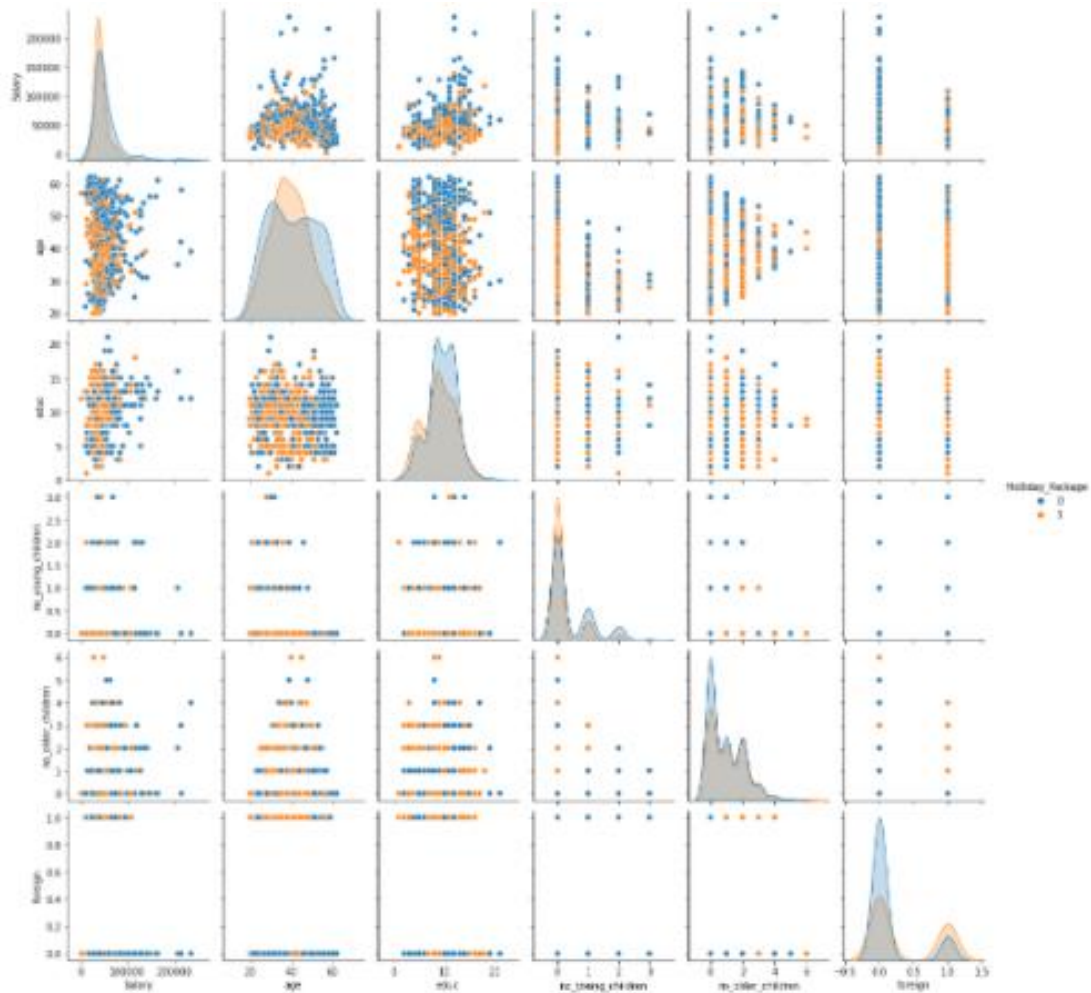


Figure 54: Pairplot

The pairplot above shows that the data points of opted and not opted overlap which means that none of the variables is a good predictor for the target column `Holiday_Package`. Foreign variable seems to be able to distinguish between opted and not opted better than most variables through this pairplot. More analysis can be withdrawn after the model building.

## Question 2.2

Do not scale the data. Encode the data (having string values) for Modelling. Data Split: Split the data into train and test (70:30). Apply Logistic Regression and LDA (linear discriminant analysis).

## Answer 2.2

Using the Label Encoder from `sklearn.preprocessing` library the string values of `Holiday_package` and `foreign` have been changed to numerical values as Logistic regression and LDA uses only numerical inputs. (Foreign: yes = 1, no = 0), (`Holiday_Package`: yes=1, no=0)

	Holliday_Package	Salary	age	educ	no_young_children	no_older_children	foreign
0	0	48412	30	8	1	1	0
1	1	37207	45	8	0	1	0
2	0	58022	46	9	0	0	0
3	0	66503	31	11	2	0	0
4	0	66734	44	12	0	2	0

Figure 55: Label Encoding

The data has been splitted into two parts which is one has the independent variables and other has the target variable(Holliday\_Package).

Using train\_test\_split function the data has been splitted into 70% train and 30% test.

Logistic Regression has been applied to the train data to train the model and further the same models are used to predict on the test data.

Logistic Regression

A Grid Search CV has been used to find out the best parameters for logistic regression

```
GridSearchCV(cv=3, estimator=LogisticRegression(max_iter=100000, n_jobs=2),
             n_jobs=-1,
             param_grid={'penalty': ['l1', 'l2', 'elasticnet', 'none'],
                          'solver': ['newton-cg', 'lbfgs', 'liblinear', 'sag',
                                      'saga'],
                          'tol': [0.001, 0.0001]},
             scoring='f1')
```

Figure 56:Grid Search CV

The best parameters that Grid Search CV keeping the scoring parameter as F1 score gives are in the image below:

```
{'penalty': 'l2', 'solver': 'newton-cg', 'tol': 0.001}
LogisticRegression(max_iter=100000, n_jobs=2, solver='newton-cg', tol=0.001)
```

Figure 57: Best Parameters

Linear Discriminant Analysis:

Similarly, like logistic regression data has been splitted to independent variables and target variable.

After this the data has been splitted into 70% train and 30% test data. Linear Discriminant Analysis has been applied to the train data to train the model and further the same models are used to predict on the test data.

```
clf=LinearDiscriminantAnalysis()
model=clf.fit(X_train,Y_train)
model

LinearDiscriminantAnalysis()
```

Figure 58: Linear Discriminant Analysis

### Question 2.3

Performance Metrics: Check the performance of Predictions on Train and Test sets using Accuracy, Confusion Matrix, Plot ROC curve and get ROC\_AUC score for each model Final Model: Compare Both the models and write inference which model is best/optimized.

### Answer 2.3

**Models without treating any outliers:**

**Logistic Regression:**

Train data

The Area under the curve is 0.742 for training dataset. Higher the AOC value better is the model so let's understand all other performance metrics.




Figure 59: ROC Curve

The accuracy is 68% but the recall for 1 is average. The parameters for 1 are more important because it tells us about the employees that have opted for the holiday package.

	precision	recall	f1-score	support
0	0.68	0.77	0.72	326
1	0.69	0.57	0.63	284
accuracy			0.68	610
macro avg	0.68	0.67	0.67	610
weighted avg	0.68	0.68	0.68	610

Figure 60: Classification Report

Confusion matrix cells are populated by the terms:

True Positive(TP)- The values which are predicted as True and are actually True.

True Negative(TN)- The values which are predicted as False and are actually False.

False Positive(FP)- The values which are predicted as True but are actually False.

False Negative(FN)- The values which are predicted as False but are actually True.

The False negatives in this case is high which is the reason for a low recall score of 1.

163 records are the ones predicted correctly for employees who have opted. 121 records are the employees who had opted but the model has predicted it wrong which is not good. 252 records are the employees who have not opted and model also predicted them correctly. 74 records are who have not opted and model has predicted them as opted.

```
array([[252, 74],
       [121, 163]], dtype=int64)
```

Figure 61: Confusion Matrix

Test Data

The area under the curve score for train data is 0.704 which is almost in line with the training dataset.

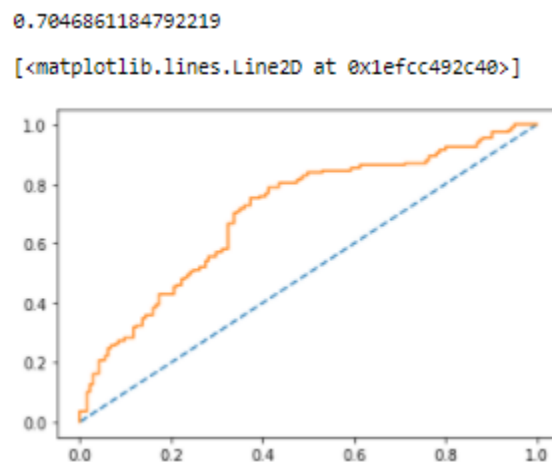


Figure 62: ROC Curve

Accuracy, recall, precision and f1 score are almost inline with the training data.

	precision	recall	f1-score	support
0	0.67	0.70	0.69	145
1	0.61	0.57	0.59	117
accuracy			0.65	262
macro avg	0.64	0.64	0.64	262
weighted avg	0.64	0.65	0.64	262

Figure 63: Classification Report

67 records are the ones predicted correctly for employees who have opted. 50 records are the employees who had opted but the model has predicted it wrong which is not good. 102 records are the employees who have not opted and model also predicted them correctly. 43 records are who have not opted and model has predicted them as opted.

```
array([[102, 43],
       [ 50, 67]], dtype=int64)
```

Figure 64: Confusion Matrix

## LDA:

Train data

The Area under the curve is 0.742 for training dataset. Higher the AOC value better is the model so let's understand all other performance metrics.

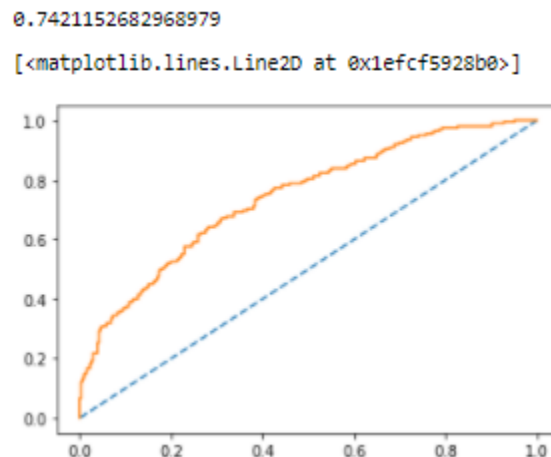


Figure 65: ROC Curve

Accuracy for training dataset is 67%. Precision and f1 score also seems to be good. Recall for 1 is little low.

	precision	recall	f1-score	support
0	0.67	0.77	0.72	326
1	0.68	0.56	0.61	284
accuracy			0.67	610
macro avg	0.67	0.66	0.66	610
weighted avg	0.67	0.67	0.67	610

Figure 66: Classification Report

158 records are the ones predicted correctly for employees who have opted. 126 records are the employees who had opted but the model has predicted it wrong which is not good. 252 records are the employees who have not opted and model also predicted them correctly. 74 records are who have not opted and model has predicted them as opted.

```
array([[252, 74],
       [126, 158]], dtype=int64)
```

Figure 67: Confusion Matrix

Test data

Area under the curve for test data is 0.702. This is inline with the train data.

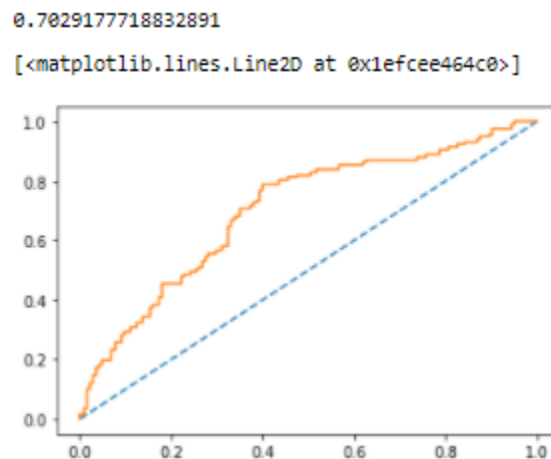


Figure 68: ROC Curve

Accuracy for test data is 64% and other parameters also are in line with the train data.

	precision	recall	f1-score	support
0	0.66	0.71	0.69	145
1	0.61	0.56	0.58	117
accuracy			0.64	262
macro avg	0.64	0.63	0.63	262
weighted avg	0.64	0.64	0.64	262

Figure 69: Classification Report



65 records are the ones predicted correctly for employees who have opted. 52 records are the employees who had opted but the model has predicted it wrong which is not good. 103 records are the employees who have not opted and model also predicted them correctly. 42 records are who have not opted and model has predicted them as opted.

```
array([[103, 42],
       [ 52, 65]], dtype=int64)
```

Figure 70: Confusion Matrix

The LDA model has been trained and tested by changing the threshold from 0.1 to 1 at an interval of 0.1. But the best results are seen at a threshold of 0.5 which is the default threshold. It gives the best combination of accuracy, f1 score, precision and recall

	Logistic Regression		LDA	
	Train	Test	Train	Test
Accuracy	0.68	0.65	0.67	0.64
AUC	0.74	0.70	0.74	0.70
F1 score	0.63	0.59	0.61	0.58
Recall	0.57	0.57	0.56	0.56
Precision	0.69	0.61	0.68	0.61

Table 2:Model comparison

The Recall, Accuracy, F1 score for logistic regression is better than LDA so the chosen model for this dataset is Logistic Regression.

### Models after outlier Treatment:

The Salary variable has a lot of outliers whereas educ, no\_young\_children, no\_older\_children are discrete variables so the outliers are only treated for the salary variable and the models are applied. The boxplot below shows that there are no more outliers left for salary variable.




Figure 71: Boxplot

### Logistic Regression Model:

A Grid Search CV is used to find the best parameters

```
GridSearchCV(cv=3, estimator=LogisticRegression(max_iter=100000, n_jobs=2),
             n_jobs=-1,
             param_grid={'penalty': ['l1', 'l2', 'elasticnet', 'none'],
                          'solver': ['newton-cg', 'lbfgs', 'liblinear', 'sag',
                                     'saga'],
                          'tol': [0.0001, 1e-05]},
             scoring='f1')
```

Figure 72: Grid Search CV

The below image gives the best parameters for Logistic Regression

```
{'penalty': 'l2', 'solver': 'newton-cg', 'tol': 0.0001}
LogisticRegression(max_iter=100000, n_jobs=2, solver='newton-cg')
```

Figure 73: Best Parameters

Train data:

Area under the curve is 0.74.




Figure 74: ROC Curve

The accuracy of the train data is 67%. Recall for 1 is little less but all other performance metrics are giving good results.

	precision	recall	f1-score	support
0	0.67	0.77	0.72	326
1	0.68	0.56	0.62	284
accuracy			0.67	610
macro avg	0.68	0.67	0.67	610
weighted avg	0.67	0.67	0.67	610

Figure 75: Classification Report

160 records are the ones predicted correctly for employees who have opted. 124 records are the employees who had opted but the model has predicted it wrong which is not good. 251 records are the employees who have not opted and model also predicted them correctly. 75 records are who have not opted and model has predicted them as opted.

```
array([[251, 75],
       [124, 160]], dtype=int64)
```

Figure 76: Confusion Matrix

Test data:

Area under the curve for test data is 0.704. This is almost in line with test data.




Figure 77: ROC Curve

The accuracy for train data is 65%. Recall has improved for test data and all metrics are in line with train data which the model is right fit.

	precision	recall	f1-score	support
0	0.67	0.70	0.69	145
1	0.61	0.57	0.59	117
accuracy			0.65	262
macro avg	0.64	0.64	0.64	262
weighted avg	0.64	0.65	0.64	262

Figure 78: Classification Report

67 records are the ones predicted correctly for employees who have opted. 50 records are the employees who had opted but the model has predicted it wrong which is not good. 102 records are the employees who have not opted and model also predicted them correctly. 43 records are who have not opted and model has predicted them as opted.

```
array([[102, 43],  
       [ 50, 67]], dtype=int64)
```

Figure 79: Confusion Matrix

## Linear Discriminant Analysis:

Train data:

Area under the curve is 0.739




Figure 80: ROC Curve

The accuracy of the train data is 68%. Recall for 1 is a little less suggesting there are more false negatives in the data but all other performance metrics look good.

	precision	recall	f1-score	support
0	0.67	0.78	0.72	326
1	0.69	0.56	0.61	284
accuracy			0.68	610
macro avg	0.68	0.67	0.67	610
weighted avg	0.68	0.68	0.67	610

Figure 81: Classification Report

158 records are the ones predicted correctly for employees who have opted. 126 records are the employees who had opted but the model has predicted it wrong which is not good. 254 records are the employees who have not opted and model also predicted them correctly. 72 records are who have not opted and model has predicted them as opted.

```
array([[254, 72],  
       [126, 158]], dtype=int64)
```

Figure 82: Confusion Matrix

Test data:

Area under the curve is 0.702




Figure 83: ROC Curve

The accuracy of the test data is 64%. Train and test data are in line with each other and model is right fit.

	precision	recall	f1-score	support
0	0.66	0.71	0.69	145
1	0.61	0.56	0.58	117
accuracy			0.64	262
macro avg	0.64	0.63	0.63	262
weighted avg	0.64	0.64	0.64	262

Figure 84: Classification Report

65 records are the ones predicted correctly for employees who have opted. 52 records are the employees who had opted but the model has predicted it wrong which is not good. 103 records are the employees who have not opted and model also predicted them correctly. 42 records are who have not opted and model has predicted them as opted.

```
array([[103, 42],  
       [ 52, 65]], dtype=int64)
```

Figure 85: ConfusionMatrix

	Logistic Regression		LDA	
	Train	Test	Train	Test
Accuracy	0.67	0.65	0.67	0.64
AUC	0.74	0.70	0.73	0.70
F1 score	0.62	0.59	0.61	0.58
Recall	0.56	0.57	0.56	0.56
Precision	0.68	0.61	0.69	0.61

Table 3:Model comparison

The better model is Logistic regression in this case as well. Accuracy, f1 score, recall is better for Logistic Regression and it is a right fit model.

## Question 2.4

Inference: Basis on these predictions, what are the insights and recommendations.

## Answer 2.4

Logistic Regression is the better model so checking the coefficients for each factor.

```
array([[ -1.74320515e-05, -5.29508007e-02,  7.15166689e-02,  
        -1.45899971e+00, -4.63494204e-02,  1.47629800e+00]])
```

Figure 86: Coefficients

1. Coefficient values for Salary is: -1.74320515e-05
2. Coefficient values for age is: -5.29508007e-02
3. Coefficient values for is educ: 7.15166689e-02
4. Coefficient values for is no\_young\_children: -1.45899971e+00
5. Coefficient values for no\_older\_children: -4.63494204e-02
6. Coefficient values for foreign: 1.47629800e+00

The most important factors affecting the target variable Holliday\_Package are no\_young\_children and foreign. Salary seemed to be one of the important factors but after model building Salary does not seem to affect the target variable.

No\_young\_children and foreign have emerged out to be strong predictors. Salary, age, educ and no\_older\_children are bad predictors.

No\_young\_children have negative coefficient which means that more the number of young children employee has it is more unlikely for him to opt for the package.

Foreign has a positive coefficient meaning more foreign employees are opting for the package.

Recommendations:

1. Company to should focus on foreign employees to drive more sales.
2. Employees with young children do not seem to opt for the package so the company can come up with a package for such employees where they can take their children also for the holiday. Although employees with young children avoid going to trips so the company should not focus more on them but can target employees who do not have any young children.

3. Company can plan some marketing and better offers to convert more employees to opt for holiday package.
4. They can offer some discounts to employees with less salary.